Medical Compute

using ChRIS on the MOC

PowerPC & x86_64 GPU Usage & Benchmarking

Bowen Jia | Elizabeth Slade | Haoyang Wang
Shineun Yoon | Kefan Zhang

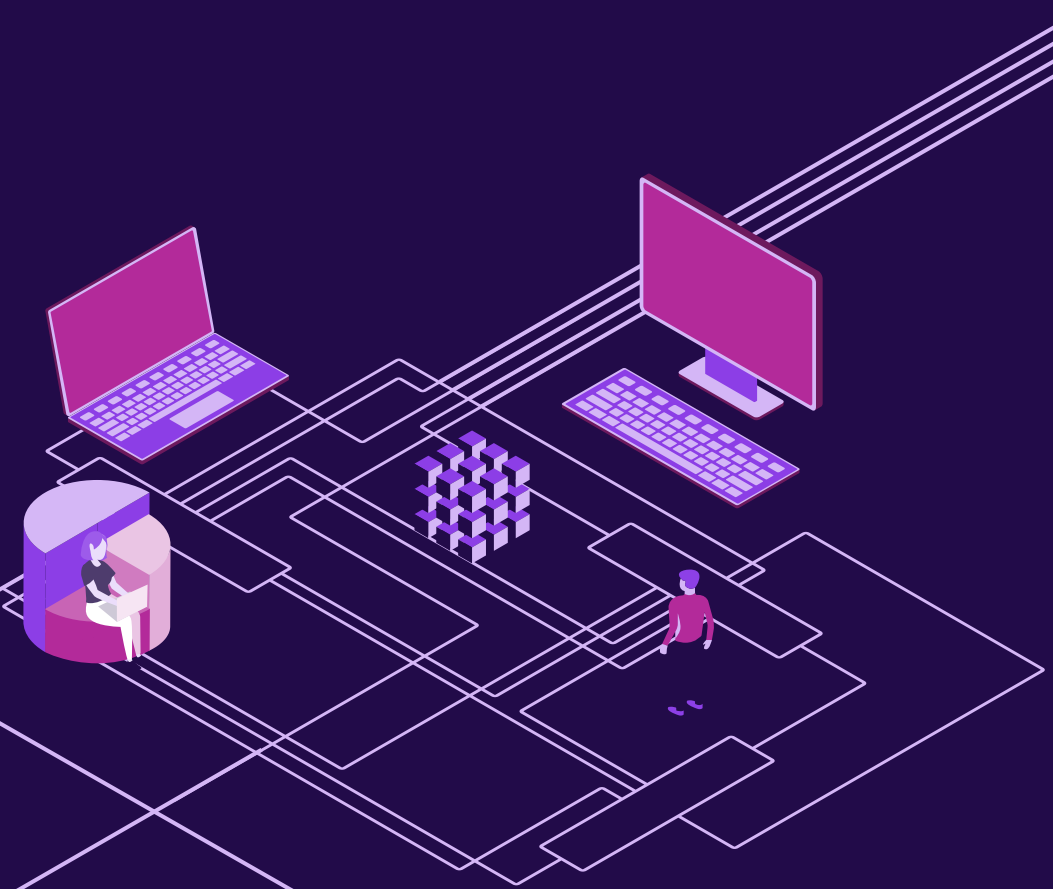

TABLE OF CONTENTS

01

Project Overview

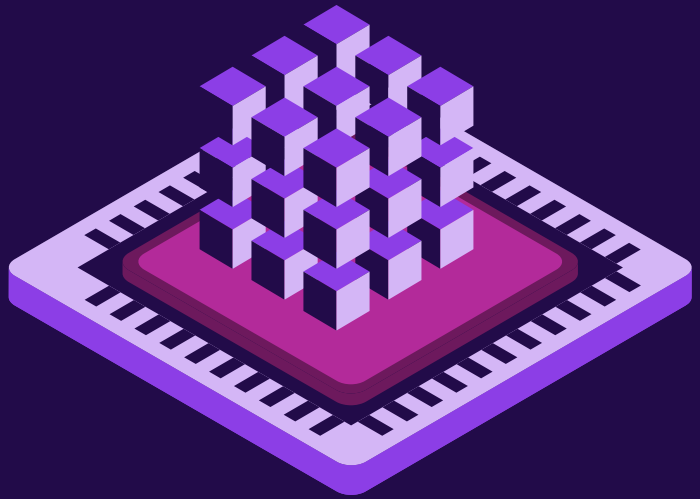

What is ChRIS?

An open source platform for medical analysis

Why ChRIS?

Technology in the medical industry is very *outdated*.
So, the Boston Children's Hospital team developed *ChRIS* to make *advanced technologies* like cloud computing and image analysis *more accessible* to doctors and researchers.

ChRIS ARCHITECTURE

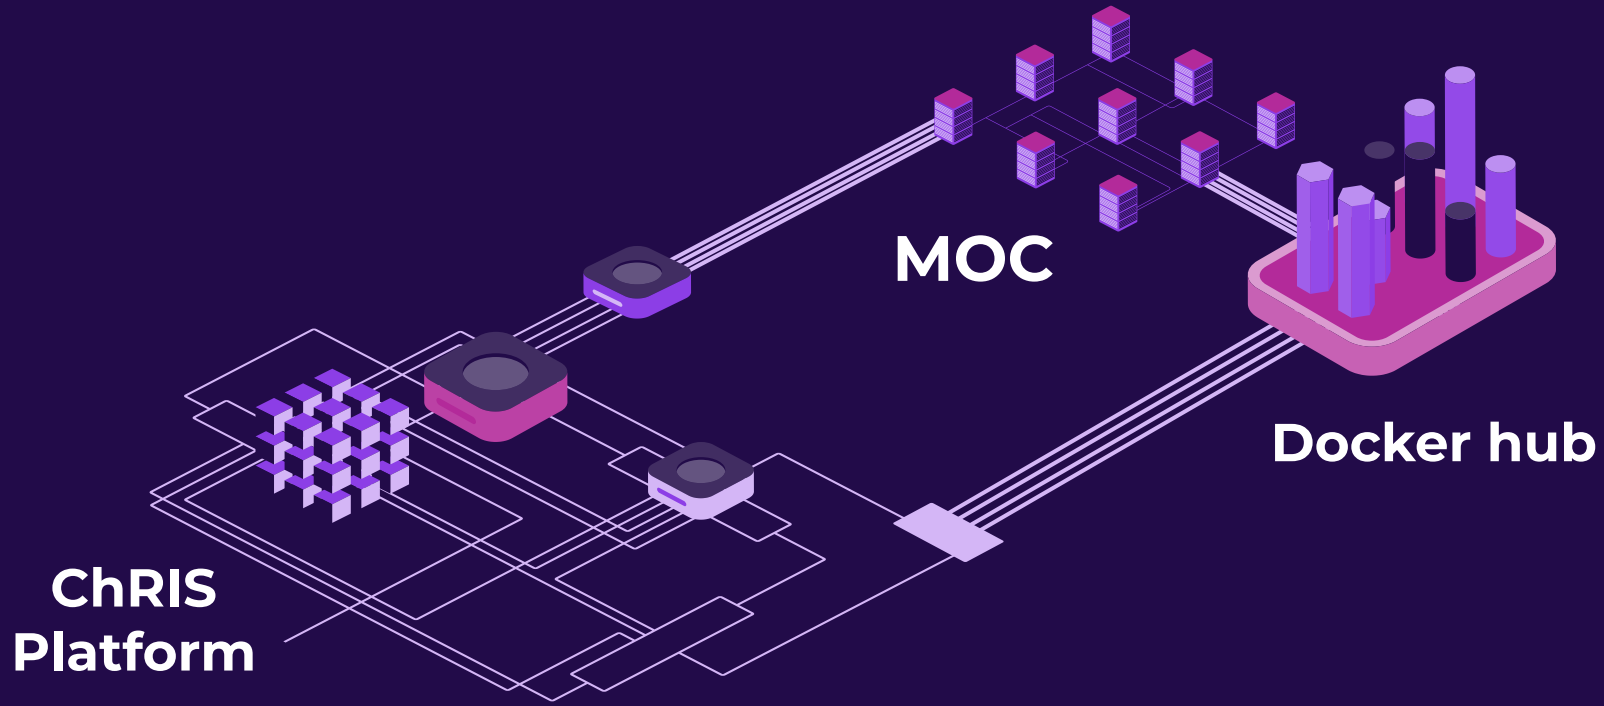

USER STORY

As a ChRIS developer / administrator,
I would like to have a way to test how my plugin performs on different architectures such as x86_64 vs PowerPC therefore I want a ChRIS plugin that performs benchmarking tests on these architectures

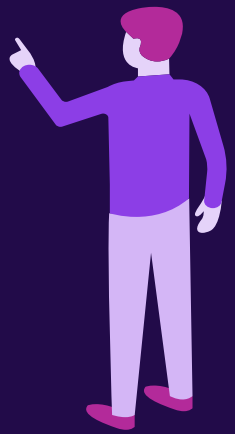

OUR PROJECT

Correctly develop
a runnable ChRIS plugin
that presents the
performance differences
between different platform
architectures

PROJECT GOALS

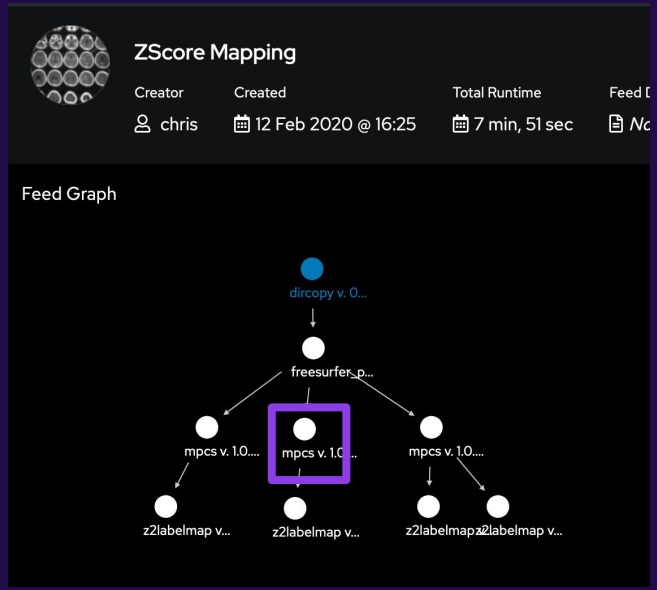

SCOPE & FEATURES

The Benchmarking Plugins should...

- Provide a series of tools to test the performance of the system
- Represent real workloads that may be deployed on the system
- Test functions should be small-sized and estimate the time that may be spent on running real computing tasks

02

Project Accomplishments

ACCOMPLISHMENTS

Docker Image 1	Docker Image 2	Docker Image 3	Docker Image 4
			
X86 Matrix Multiply Plug-in	PowerPC Matrix Multiply Plug-in	X86 Object Detection Plug-in	PowerPC Object Detection Plug-in

PowerPC pman and pfioh images on MOC are available to Boston Children's team

ChRIS WORKFLOW

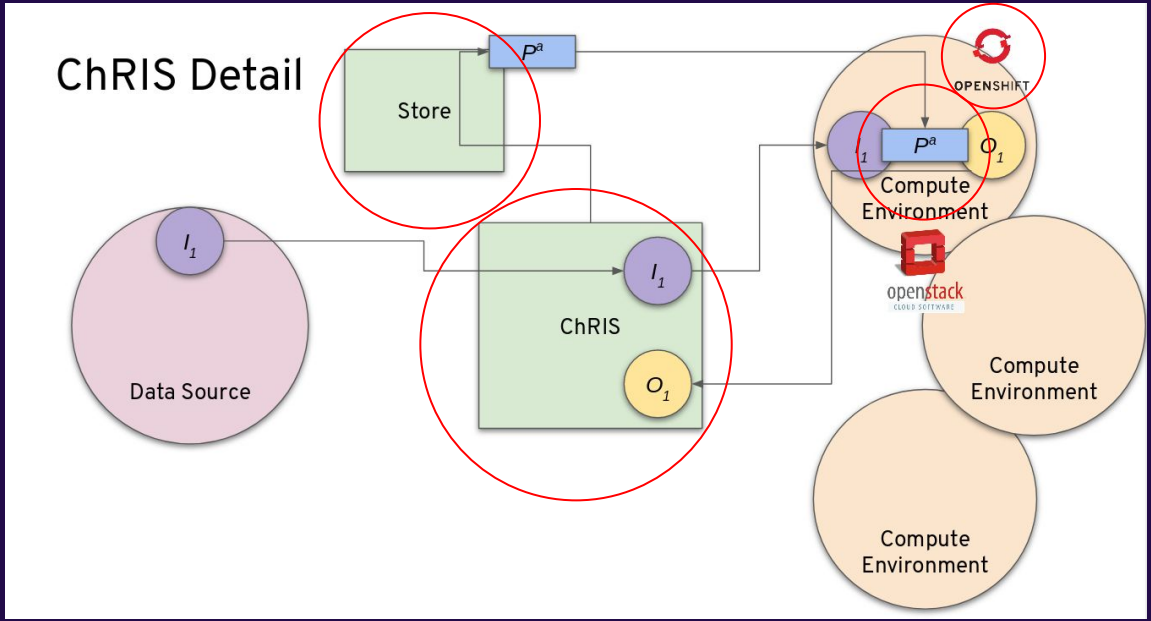

INDIVIDUAL APPS

Matrix Multiplication

Object Detection

MATRIX MULTIPLICATION

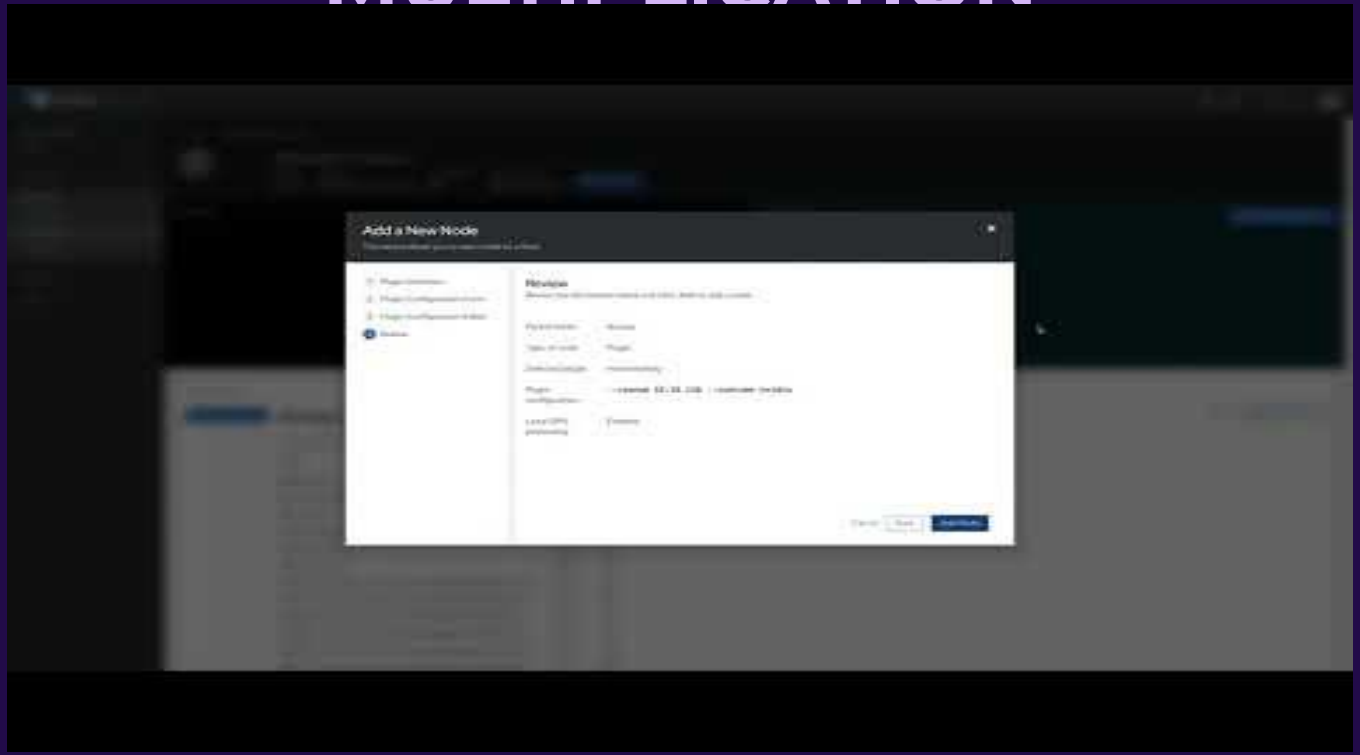

OBJECT DETECTION

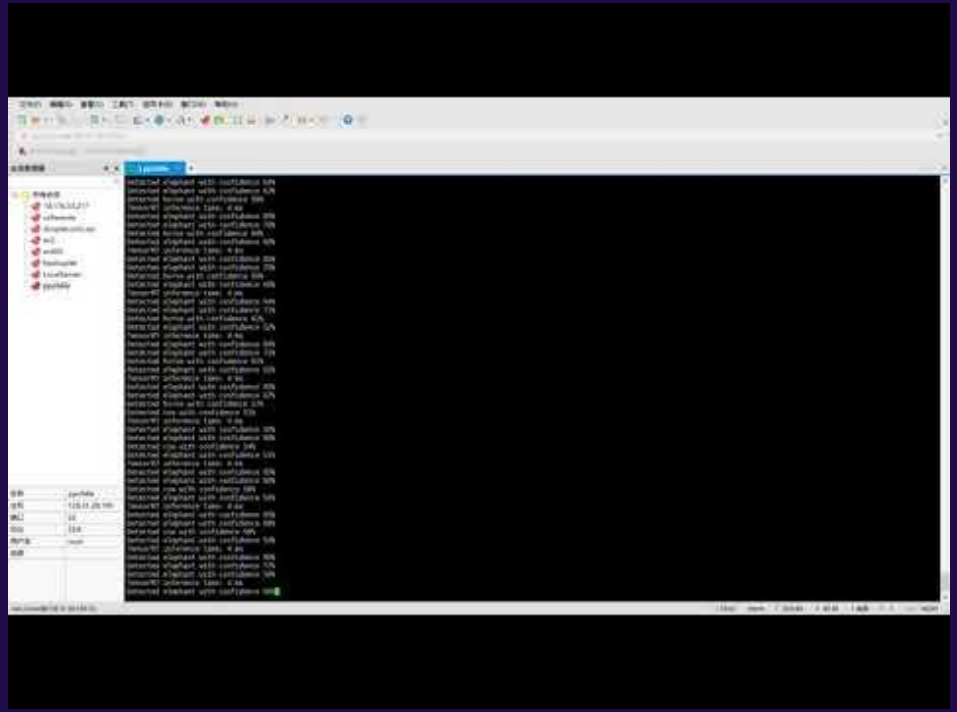

BENCHMARKING RESULT

- Matrix Multiplication:

x86:

```
(chrisapp_env) [chris-local@titan:x86_64-Linux]..._ultron_backend(master@9515-0019-)$>more MultiplyRecord.csv-1
Matrix_Size,Start_Time,Finish_Time,Elapse_Time
1048576.0,1586980604.8279586,1586980605.4869623,0.659003734588623
4194304.0,1586980605.4876544,1586980605.6453953,0.15774083137512207
9437184.0,1586980605.6458883,1586980606.0912654,0.4453771114349365
16777216.0,1586980606.0923328,1586980607.0525713,0.9602384567260742
```

Power9:

```
[root@bu-21-9 out]# cat MultiplyRecord.csv
Matrix_Size,Start_Time,Finish_Time,Elapse_Time
1048576.0,1588562635.3954642,1588562636.0703576,0.6748933792114258
4194304.0,1588562636.0713058,1588562636.1895485,0.11824274063110352
9437184.0,1588562636.1904578,1588562636.5205216,0.3300638198852539
16777216.0,1588562636.522102,1588562637.2556093,0.7335071563720703
```

BENCHMARKING RESULT

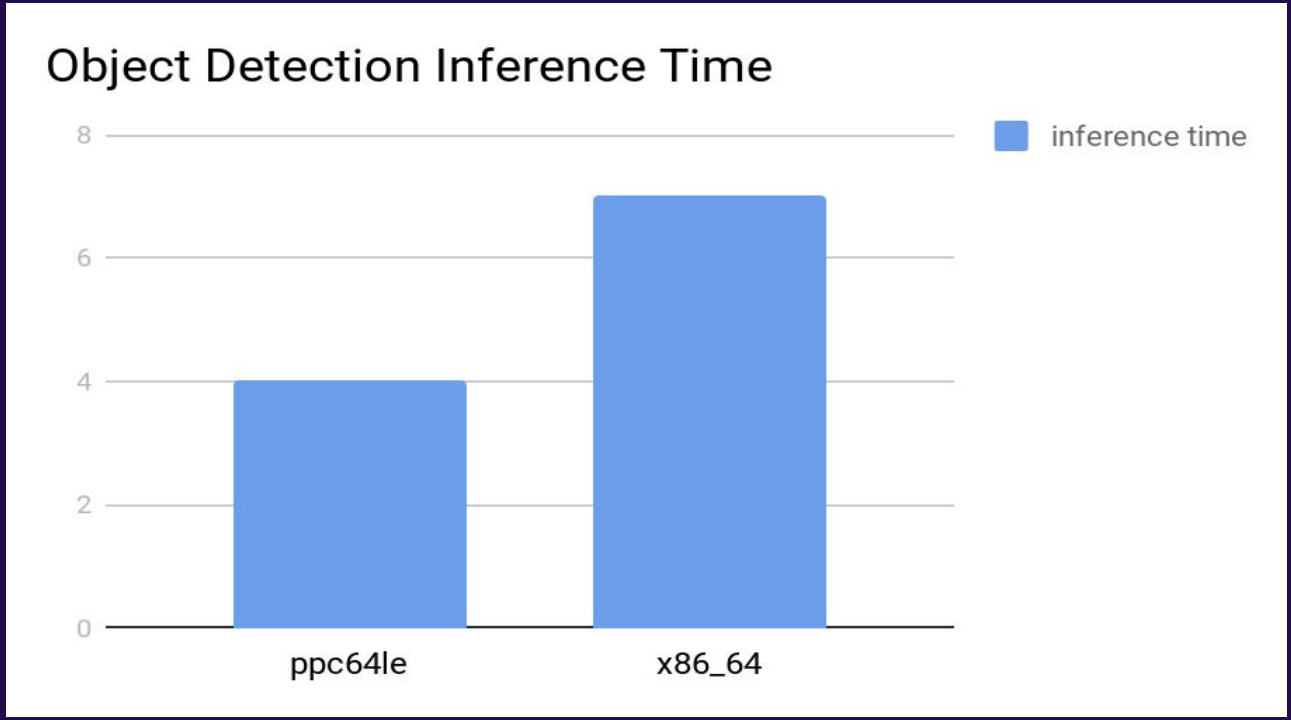

MILESTONES

- [x] Develop ChRIS plugins for x86 and PowerPC environments
- [x] Deploy pman & pfioh on MOC both for x86 & PowerPC cluster
- [x] Run ChRIS plugins through ChRIS GUI and CLI
- [x] Get MicroBench results using our plugins

Bonus Task:

- [ ] Run ChRIS plugins on PowerPC MOC cluster through pfcon service

Demo for MOC, ChRIS app...

Like saying hello, whatever

03

What we have learned

What we learned: Technical

- Concepts and usage of Docker and containers
- Deploying to MOC using Openshift
- Concepts of ChRIS architecture and building ChRIS apps
- GPU Programming (eg. parallel computing, GPU memory structure).
- Building PowerPC Images

Deploying PowerPC `pman` and `pfioh` on MOC

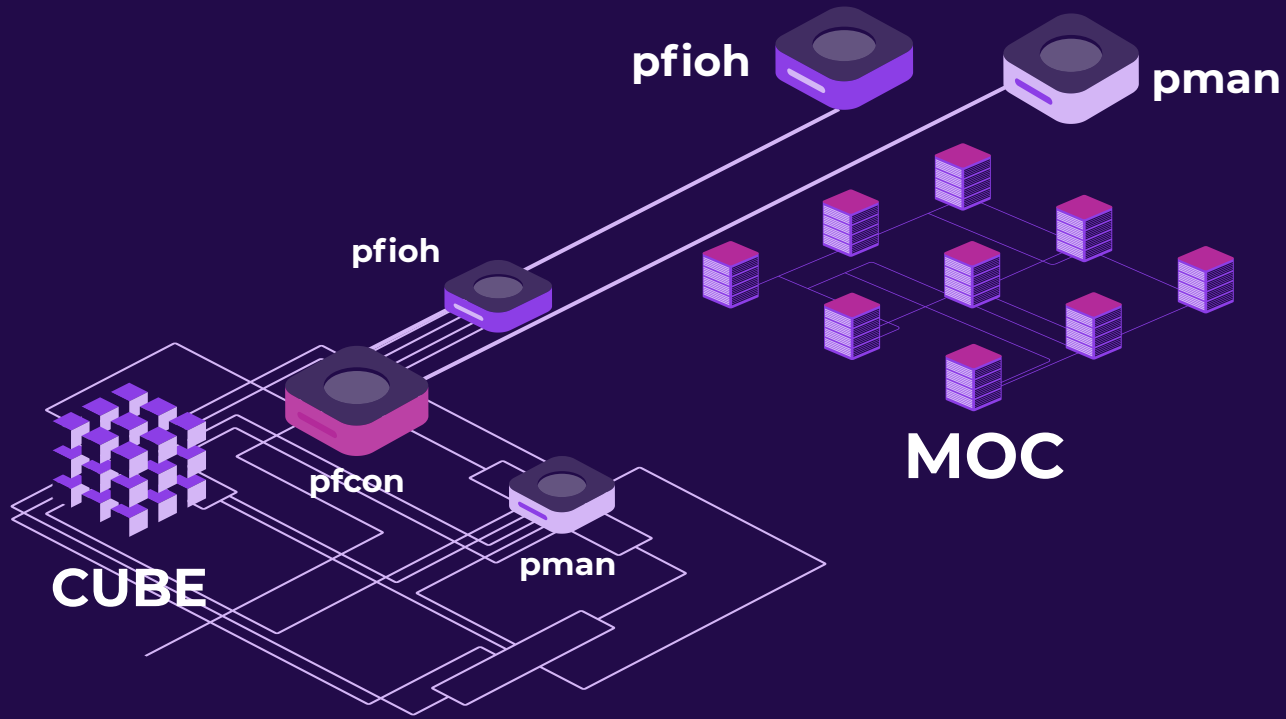

What we learned: Nontechnical

- Scrum team concept and sprints
- Managing and distributing tasks with Taiga.io
- Communicating technical challenges effectively

04

Limitations & Future Plan

LIMITATIONS

Benchmark purpose too general

Benchmark tests cannot tailor to different use cases

GPU Number Limitation

Only able to activate one GPU

FUTURE PLANS

Tool Box: More specific parameters

Parameter for focusing on optimizing iterative computations

- Machine learning plug-in, protein folding prediction

Parameter for focusing on memory read and writes speed

- File I/O, image recognition

GPU Number Limitation

Parameter for managing and processing big datasets

Brain scanning

Red Hat COVID-19

Computational projects that can help scientifically with COVID analysis

Moving project plugin to Power9 platform.

BURNDOWN CHART

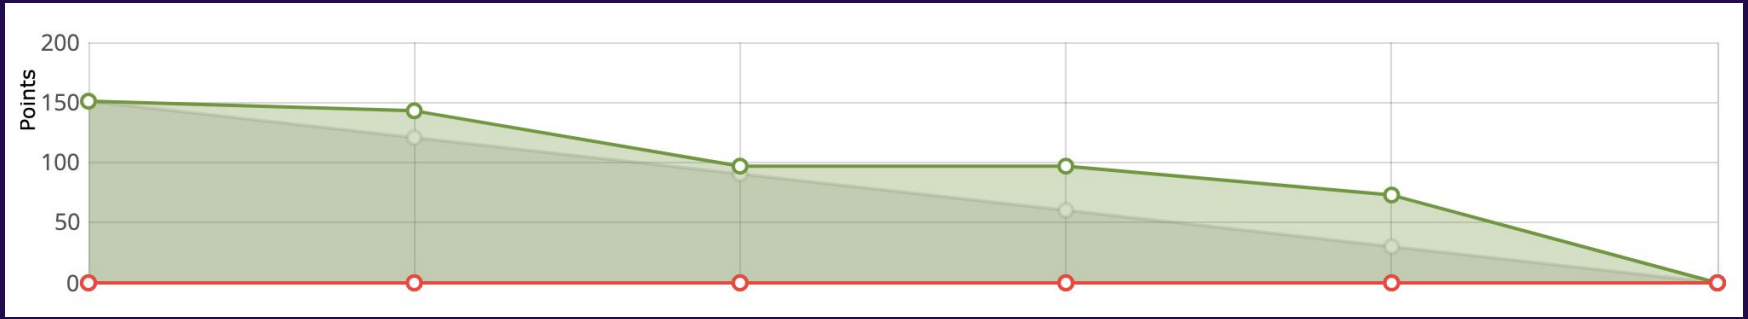

THANK YOU